

# HW4

April 14, 2020

## 1 HW4

### 1.1 Part I

Use the **PopDivas\_\_data.csv** on GitHub to build a prediction model.

1. Explore the Data
  - What patterns/relationships do you notice?
2. Choose/Build Your Model (*Decision Tree* OR *KNN*)
  - Why did you choose this type of model?
  - Which variables are you including in your model?
  - Choose a model validation technique and explain why you chose it.
  - Which variables did you standardize and Why?
3. Evaluate Your Model
  - State how it did (and what evidence/metric do you have to support that?)

### 1.2 Part II

Use the **YouTubeKidsVideo.csv** on GitHub to build a Naive Bayes Classifier. This dataset looks at the titles/descriptions of YouTube videos that are (1) and are not (0) meant for kids. The variable KidsVideo is 1 if the video is meant for kids, and 0 if it is not. The other variables are 1 if that word (e.g. “toy”, “girl”...etc) is in the title/description of the video, and 0 if it is not.

1. Explore the Data
  - What patterns/relationships do you notice?
2. Build your model
  - Which variables are you including in your model?
  - Choose a model validation technique and explain why you chose it.
3. Evaluate Your Model
  - How did it do? What evidence/metric do you have to support that?

## 2 Part I

```
[87]: import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
```

```

from plotnine import *
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier, KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB

from sklearn import metrics
from sklearn.preprocessing import StandardScaler, LabelBinarizer
from sklearn.model_selection import KFold # k-fold cv
from sklearn.model_selection import cross_val_score # cross validation metrics
from sklearn.model_selection import cross_val_predict # cross validation metrics
from sklearn.metrics import accuracy_score, confusion_matrix, mean_squared_error
#from sklearn.metrics import plot_confusion_matrix

from sklearn.model_selection import GridSearchCV

%precision %.7g
%matplotlib inline

```

```

[88]: divas = pd.read_csv('data/PopDivas_data.csv')
print(divas.head())
print(divas.info())
print(divas.isnull().sum())
print(divas.describe())

```

```

      Unnamed: 0  artist_name  danceability  energy  key  loudness  mode  \
0              1    Beyoncé      0.386  0.28800   1    -18.513    1
1              2    Beyoncé      0.484  0.36300   5     -8.094    0
2              3    Beyoncé      0.537  0.24700   2    -17.750    1
3              4    Beyoncé      0.672  0.69600   4     -6.693    0
4              5    Beyoncé      0.000  0.00515   9    -22.612    0

      speechiness  acousticness  instrumentalness  liveness  valence  \
0          0.0602          0.533          0.01670    0.1410    0.399
1          0.0368          0.645          0.00000    0.1250    0.201
2          0.0793          0.199          0.00001    0.4230    0.170
3          0.1770          0.200          0.02750    0.0736    0.642
4          0.0000          0.524          0.95000    0.1140    0.000

      duration_ms  track_name
0          43850  balance (mufasa interlude)
1          226479                BIGGER
2          46566  the stars (mufasa interlude)
3          162353        FIND YOUR WAY BACK
4          13853  uncle scar (scar interlude)
<class 'pandas.core.frame.DataFrame'>

```

RangeIndex: 1599 entries, 0 to 1598

Data columns (total 14 columns):

Unnamed: 0            1599 non-null int64  
artist\_name           1599 non-null object  
danceability          1599 non-null float64  
energy                1599 non-null float64  
key                   1599 non-null int64  
loudness              1599 non-null float64  
mode                  1599 non-null int64  
speechiness           1599 non-null float64  
acousticness          1599 non-null float64  
instrumentalness      1599 non-null float64  
liveness              1599 non-null float64  
valence               1599 non-null float64  
duration\_ms           1599 non-null int64  
track\_name            1599 non-null object  
dtypes: float64(8), int64(4), object(2)

memory usage: 175.0+ KB

None

Unnamed: 0            0  
artist\_name           0  
danceability          0  
energy                0  
key                   0  
loudness              0  
mode                  0  
speechiness           0  
acousticness          0  
instrumentalness      0  
liveness              0  
valence               0  
duration\_ms           0  
track\_name            0

dtype: int64

	Unnamed: 0	danceability	energy	key	loudness \
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	800.000000	0.602417	0.645139	5.329581	-6.59192
std	461.735855	0.161994	0.199446	3.506414	3.25770
min	1.000000	0.000000	0.005150	0.000000	-27.63100
25%	400.500000	0.494000	0.517500	2.000000	-7.33900
50%	800.000000	0.628000	0.684000	6.000000	-5.82500
75%	1199.500000	0.718000	0.797000	8.000000	-4.76400
max	1599.000000	0.932000	0.994000	11.000000	-1.35700

	mode	speechiness	acousticness	instrumentalness	liveness \
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.602877	0.120019	0.205701	0.015319	0.232994
std	0.489455	0.135377	0.261146	0.094224	0.208165

min	0.000000	0.000000	0.000025	0.000000	0.016200
25%	0.000000	0.040100	0.016200	0.000000	0.097800
50%	1.000000	0.061500	0.082500	0.000002	0.149000
75%	1.000000	0.151000	0.302500	0.000081	0.310000
max	1.000000	0.960000	0.977000	0.955000	0.993000

	valence	duration_ms
count	1599.000000	1.599000e+03
mean	0.480844	2.215496e+05
std	0.233011	6.732088e+04
min	0.000000	8.619000e+03
25%	0.293000	1.997860e+05
50%	0.467000	2.229860e+05
75%	0.664000	2.491995e+05
max	0.971000	1.187253e+06

```
[89]: print(divas['artist_name'].unique())
      print(divas['key'].unique())
```

```
['Beyoncé' 'Britney Spears' 'Christina Aguilera' 'Lady Gaga' 'Rihanna'
 'Ariana Grande']
[ 1  5  2  4  9  7  6  0  8  3 10 11]
```

```
[90]: divas['artist_name'] = divas['artist_name'].astype('category')

label_binary = LabelBinarizer()
lb_results = label_binary.fit_transform(divas["artist_name"])
column_names = label_binary.classes_
# lb_results.shape
# column_names

lb_df = pd.DataFrame(lb_results, columns = column_names)
#print(lb_df)

divas = divas.join(lb_df, lsuffix='index', rsuffix='index')

print(divas.head(20))
#combined.drop(combined.columns[9], axis=1, inplace=True)
column_names = divas.columns[15:]
```

	Unnamed: 0	artist_name	danceability	energy	key	loudness	mode	\
0	1	Beyoncé	0.386	0.28800	1	-18.513	1	
1	2	Beyoncé	0.484	0.36300	5	-8.094	0	
2	3	Beyoncé	0.537	0.24700	2	-17.750	1	
3	4	Beyoncé	0.672	0.69600	4	-6.693	0	
4	5	Beyoncé	0.000	0.00515	9	-22.612	0	
5	6	Beyoncé	0.932	0.77500	7	-5.345	0	
6	7	Beyoncé	0.563	0.47000	6	-13.470	1	

7	8	Beyoncé	0.790	0.69300	0	-5.767	1
8	9	Beyoncé	0.489	0.35300	8	-17.619	1
9	10	Beyoncé	0.484	0.45700	8	-9.458	0
10	11	Beyoncé	0.494	0.41100	6	-15.021	0
11	12	Beyoncé	0.608	0.70900	7	-6.175	1
12	13	Beyoncé	0.000	0.55800	0	-17.238	1
13	14	Beyoncé	0.804	0.82300	0	-3.667	0
14	15	Beyoncé	0.603	0.60200	6	-7.083	1
15	16	Beyoncé	0.000	0.18600	0	-17.370	1
16	17	Beyoncé	0.731	0.67400	2	-6.704	1
17	18	Beyoncé	0.628	0.28700	9	-16.666	0
18	19	Beyoncé	0.647	0.73200	2	-5.846	1
19	20	Beyoncé	0.255	0.32000	8	-19.821	1

	speechiness	acousticness	instrumentalness	liveness	valence	\
0	0.0602	0.5330	0.016700	0.1410	0.399	
1	0.0368	0.6450	0.000000	0.1250	0.201	
2	0.0793	0.1990	0.000010	0.4230	0.170	
3	0.1770	0.2000	0.027500	0.0736	0.642	
4	0.0000	0.5240	0.950000	0.1140	0.000	
5	0.1150	0.0184	0.015700	0.3180	0.584	
6	0.1470	0.0306	0.661000	0.2830	0.186	
7	0.0605	0.0420	0.026700	0.0790	0.855	
8	0.6640	0.4910	0.000022	0.2290	0.175	
9	0.2040	0.7380	0.000000	0.5360	0.342	
10	0.7710	0.7030	0.000000	0.1560	0.403	
11	0.3740	0.1190	0.000000	0.3370	0.709	
12	0.0000	0.7560	0.000000	0.5330	0.000	
13	0.0745	0.0984	0.000004	0.2460	0.489	
14	0.2300	0.0490	0.000000	0.2790	0.595	
15	0.0000	0.7830	0.000140	0.3330	0.000	
16	0.0801	0.0057	0.000025	0.1090	0.730	
17	0.6600	0.6750	0.000000	0.4260	0.535	
18	0.0891	0.0846	0.000000	0.5970	0.548	
19	0.0533	0.8260	0.021900	0.2610	0.302	

	duration_ms	track_name	\
0	43850	balance (mufasa interlude)	
1	226479	BIGGER	
2	46566	the stars (mufasa interlude)	
3	162353	FIND YOUR WAY BACK	
4	13853	uncle scar (scar interlude)	
5	155990	DON'T JEALOUS ME	
6	16768	danger (young simba & young nala interlude)	
7	190108	JA ARA E	
8	29582	run away (scar & young simba interlude)	
9	107204	NILE	
10	53793	new lesson (timon, pumbaa & young simba interl...	

11	272068	MOOD 4 EVA (feat. Oumou Sangaré)
12	8619	reunited (nala & simba interlude)
13	152563	WATER
14	248472	BROWN SKIN GIRL
15	14827	come home (nala interlude)
16	198952	KEYS TO THE KINGDOM
17	31391	follow me (simba & rafiki interlude)
18	222529	ALREADY
19	45293	remember (mufasa interlude)

	Ariana Grande	Beyoncé	Britney Spears	Christina Aguilera	Lady Gaga	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	
5	0	1	0	0	0	
6	0	1	0	0	0	
7	0	1	0	0	0	
8	0	1	0	0	0	
9	0	1	0	0	0	
10	0	1	0	0	0	
11	0	1	0	0	0	
12	0	1	0	0	0	
13	0	1	0	0	0	
14	0	1	0	0	0	
15	0	1	0	0	0	
16	0	1	0	0	0	
17	0	1	0	0	0	
18	0	1	0	0	0	
19	0	1	0	0	0	

	Rihanna
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	0
8	0
9	0
10	0
11	0
12	0
13	0
14	0

```

15      0
16      0
17      0
18      0
19      0

```

```
[91]: divas.head()
```

```

[91]:   Unnamed: 0  artist_name  danceability  energy  key  loudness  mode  \
0          1    Beyoncé      0.386  0.28800    1   -18.513    1
1          2    Beyoncé      0.484  0.36300    5    -8.094    0
2          3    Beyoncé      0.537  0.24700    2   -17.750    1
3          4    Beyoncé      0.672  0.69600    4    -6.693    0
4          5    Beyoncé      0.000  0.00515    9   -22.612    0

      speechiness  acousticness  instrumentalness  liveness  valence  \
0      0.0602      0.533      0.01670    0.1410    0.399
1      0.0368      0.645      0.00000    0.1250    0.201
2      0.0793      0.199      0.00001    0.4230    0.170
3      0.1770      0.200      0.02750    0.0736    0.642
4      0.0000      0.524      0.95000    0.1140    0.000

      duration_ms      track_name  Ariana Grande  Beyoncé  \
0      43850  balance (mufasa interlude)          0        1
1      226479                BIGGER          0        1
2      46566  the stars (mufasa interlude)          0        1
3      162353      FIND YOUR WAY BACK          0        1
4      13853  uncle scar (scar interlude)          0        1


      Britney Spears  Christina Aguilera  Lady Gaga  Rihanna
0          0          0          0          0
1          0          0          0          0
2          0          0          0          0
3          0          0          0          0
4          0          0          0          0

```

```

[92]: (ggplot(divas, aes(x = "energy", y = "danceability"))
      +geom_point())

```




[92]: <ggplot: (-9223371894412295976)>

1a - Seems like there is generally a positive relationship between energy and dancability due to the large cluster in the top right corner

```
[93]: (ggplot(divas, aes(x = "duration_ms", y = "danceability"))  
      +geom_point())
```





```
[93]: <ggplot: (-9223371894413343008)>
```

data is too bunched up to view in this view

```
[94]: divas['duration_s'] = divas['duration_ms'] / 1000

# get rid of interlude songs
divas = divas[divas['duration_s'] > 60]
print(divas.shape)
```

```
(1539, 21)
```

```
[95]: #(divas['duration_s'] > 600).sum()
print(divas.loc[divas['duration_s'] > 600])

# got rid of that mix because it would likely throw off the model
divas = divas[divas['duration_s'] < 600]
```

```

      Unnamed: 0  artist_name  danceability  energy  key  loudness  mode  \
268          269    Beyoncé         0.262    0.971   11    -3.827    0

      speechiness  acousticness  instrumentalness  ...  valence  duration_ms  \

```

```

268      0.383      0.108      0.0 ...      0.289      1187253


                                track_name Ariana Grande  Beyoncé \
268  Destiny's Child Medley - Audio from The Beyonc...      0      1

      Britney Spears  Christina Aguilera  Lady Gaga  Rihanna  duration_s
268              0              0              0              0      1187.253

[1 rows x 21 columns]

```

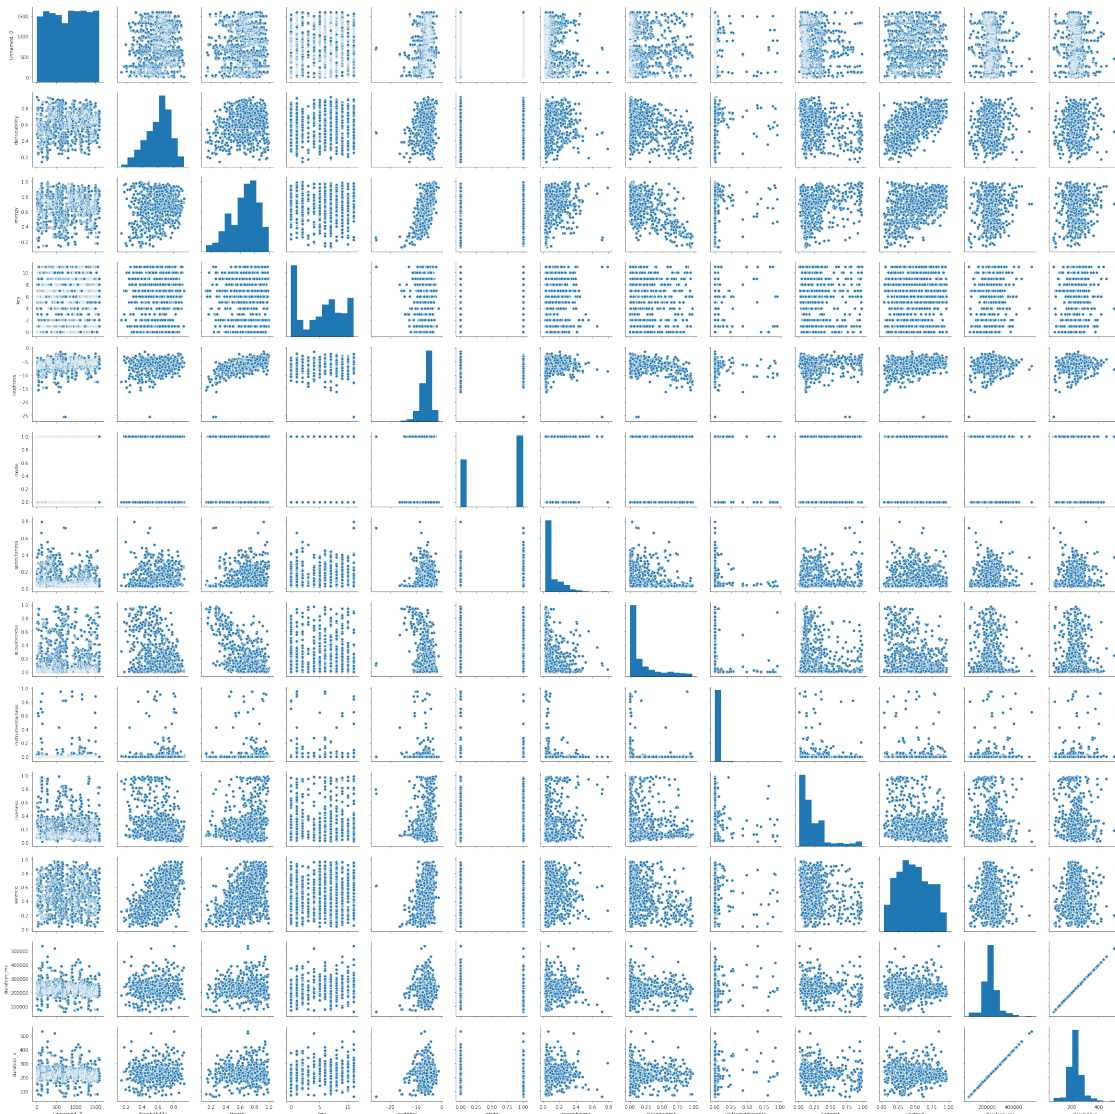
```
[8]: (ggplot(divas, aes(x = "duration_s", y = "danceability"))
      +geom_point())
```



```
[8]: <ggplot: (-9223371894318240052)>
```

1b - converted to seconds so that its easier to interpret. also removed the destinys child medley so that it wouldnt throw off the model

```
[16]: g = sns.pairplot(divas)
```



1c - at a glance valence also looks like it has a positive relationship, as well as loudness a little bit.

```
[113]: predictors = ["valence", "duration_s", "loudness"]

predictors[2:2] = column_names

#print(predictors)
X = divas[predictors]
y = divas["danceability"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)

knn = KNeighborsRegressor()
```

```

ks = {"n_neighbors": range(1,30)}

# use grid search to find best parameters
grid = GridSearchCV(knn,ks, scoring = "r2", cv = 5)

zscore = StandardScaler()
zscore.fit(X_train)

Xs_train = zscore.transform(X_train)
Xs_test = zscore.transform(X_test)

knnmod = grid.fit(X_train, y_train)

knnmod.best_estimator_.get_params()["n_neighbors"]

```

[113]: 1

```

[114]: poss_k = [1,2,3,4,5,6,7,8,9,10]
acc = {}

for k in poss_k:
    kf = KFold(n_splits = 5)
    knn2 = KNeighborsRegressor(n_neighbors = k)

    acc[k] = np.mean(cross_val_score(knn2, X_train, y_train, cv = kf))

print(acc)

chosen_k = max(acc, key=acc.get)
print(chosen_k)

knn_final = KNeighborsRegressor(n_neighbors = chosen_k)
knn_final.fit(X_train,y_train)

knn_final.score(X_test,y_test)

```

```

{1: 0.3247511583538543, 2: 0.30816326432983165, 3: 0.2830303003513635, 4:
0.25626460996674727, 5: 0.23821537517322752, 6: 0.20528739606340496, 7:
0.18350167613503354, 8: 0.15760007469868254, 9: 0.14593655363540706, 10:
0.1396864875637521}
1

```

[114]: 0.535546400750241

Write your responses here

2 - I chose KNN as the model to use because I believe artists draw a lot of inspiration from other

artists and there might be an interesting relationship between these artists since they were kind of around at the same time and I think knn is an interesting model that might be able to capture this relationship.

```
[109]: train_pred = knnmod.predict(X_train)
        test_pred = knnmod.predict(X_test)

        print('training r2 is:', knnmod.score(X_train, y_train)) #training R2
        print('testing r2 is:', knnmod.score(X_test, y_test)) #testing R2

        print('\ntrain mse is: ', mean_squared_error(y_train, train_pred))
        print('test mse is: ', mean_squared_error(y_test, test_pred))
```

```
training r2 is: 0.9999997641015631
testing r2 is: 0.19972686224014524
```

```
train mse is: 5.691056910568755e-09
test mse is: 0.017059165584415583
```

```
[115]: train_pred = knn_final.predict(X_train)
        test_pred = knn_final.predict(X_test)

        print('training r2 is:', knn_final.score(X_train, y_train)) #training R2
        print('testing r2 is:', knn_final.score(X_test, y_test)) #testing R2

        print('\ntrain mse is: ', mean_squared_error(y_train, train_pred))
        print('test mse is: ', mean_squared_error(y_test, test_pred))
```

```
training r2 is: 0.9999997915786427
testing r2 is: 0.535546400750241
```

```
train mse is: 4.878048780487634e-09
test mse is: 0.011231870129870132
```

Write your responses here

The model didnt do too terrible according to the r2 in the k-fold model. The initial model seemed very overfitted to the training data which is why i decided to make a k-fold model as well. I standardized all variables because I couldnt figure out how to choose specific ones. I kept getting errors on the zscore fit method.

### 3 Part II

Use the YouTubeKidsVideo.csv on GitHub to build a Naive Bayes Classifier. This dataset looks at the titles/descriptions of YouTube videos that are (1) and are not (0) meant for kids. The variable KidsVideo is 1 if the video is meant for kids, and 0 if it is not. The other variables are 1 if that word (e.g. “toy”, “girl”...etc) is in the title/description of the video, and 0 if it is not.

Explore the Data

What patterns/relationships do you notice?  
 Build your model  
 Which variables are you including in your model?  
 Choose a model validation technique and explain why you chose it.  
 Evaluate Your Model  
 How did it do? What evidence/metric do you have to support that?

```
[31]: vids = pd.read_csv('data/YouTubeKidsVideo.csv')
print(vids.head())
print(vids.info())
print(vids.isnull().sum())
print(vids.describe())
```

```
   kidsVideo  cat  toy  sad  girl  is
0          1    1    0    0     0    0
1          1    1    0    0     0    0
2          1    1    1    0     0    0
3          1    1    1    0     0    0
4          1    0    0    0     0    0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 6 columns):
kidsVideo    2000 non-null int64
cat           2000 non-null int64
toy           2000 non-null int64
sad           2000 non-null int64
girl         2000 non-null int64
is            2000 non-null int64
dtypes: int64(6)
memory usage: 93.9 KB
None
kidsVideo    0
cat           0
toy           0
sad           0
girl          0
is            0
dtype: int64
```

	kidsVideo	cat	toy	sad	girl \
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	0.500000	0.488000	0.304500	0.234000	0.236000
std	0.500125	0.499981	0.460310	0.423478	0.424728
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.500000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	is
count	2000.000000
mean	0.145000
std	0.352189
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000


```
[32]: kids = vids[vids['kidsVideo'] == 1]
      not_kids = vids[vids['kidsVideo'] == 0]

      print(kids.shape)
      print(not_kids.shape)
```

```
(1000, 6)
(1000, 6)
```


```
[33]: Data = {'total': [kids['cat'].sum(),kids['toy'].sum(),kids['sad'].
      ↪sum(),kids['girl'].sum(),kids['is'].sum()]}
      df = pd.DataFrame(Data,columns=['total'])
      my_labels = 'cat','toy','sad','girl','is'
      plt.pie(df,labels=my_labels,autopct='%1.1f%%')
      plt.title('Word Prevalance in Title for YouTube Videos MEANT for Kids')
      plt.axis('equal')
      plt.show()
```

Word Prevalance in Title for YouTube Videos MEANT for Kids



```
[34]: Data = {'total': [not_kids['cat'].sum(),not_kids['toy'].sum(),not_kids['sad'].
    ↪sum(),not_kids['girl'].sum(),not_kids['is'].sum()]}
df = pd.DataFrame(Data,columns=['total'])
my_labels = 'cat','toy','sad', 'girl', 'is'
plt.pie(df,labels=my_labels,autopct='%1.1f%%')
plt.title('Word Prevalance in Title for YouTube Videos NOT MEANT for Kids')
plt.axis('equal')
plt.show()
```

Word Prevalance in Title for YouTube Videos NOT MEANT for Kids



data shows that there are more titles with the word toy in the title in videos meant for kids. The opposite is true in the case for videos not meant for children, rather sad is a more common word in titles not meant for kids.

```
[118]: predictors = ['toy', 'sad']

X = vids[predictors]
y = vids["kidsVideo"]

kf = KFold(n_splits = 4)
nb = GaussianNB()
acc = []
predictedVals = []
for train, test in kf.split(X,y):
    X_train = X.iloc[train]
```



```

X_test = X.iloc[test]
y_train = y[train]
y_test = y[test]

nb.fit(X_train,y_train)
acc.append(nb.score(X_test,y_test))
cnf_matrix = confusion_matrix(y_test, y_test)
print(cnf_matrix)

```

```

[[500]]
[[500]]
[[500]]
[[500]]

```

Write your responses here

i used the two most prevelant words in the model because they were the most distinct keywords in titles. I chose accuracy as the metric to use in assessing the model because it is a very common metric to determines the overall predicted accuracy of the model.

```

[117]: print(acc)
       print(np.mean(acc))

```

```

[0.46, 0.454, 0.68, 0.666]
0.565000000000000001

```

Write your responses here

The model doesnt do too well given a score close to .50 indicating that the model is accurate and useful in half of the time.

```

[ ]: get_ipython().system("jupyter nbconvert --output-dir='output/' --to pdf knn_nb.
    ↪ipynb")
get_ipython().system("jupyter nbconvert --output-dir='output/' --to markdown_
    ↪knn_nb.ipynb")
get_ipython().system("jupyter nbconvert --output-dir='output/' --to html knn_nb.
    ↪ipynb")
get_ipython().system("jupyter nbconvert --output-dir='output/' --to python_
    ↪knn_nb.ipynb")

```

This application is used to convert notebook files (\*.ipynb) to various other formats.

WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.

Options

-----

Arguments that take values are actually convenience aliases to full Configurables, whose aliases are listed on the help line. For more information on full configurables, see '--help-all'.

--debug

set log level to logging.DEBUG (maximize logging output)

--generate-config

generate default config file

-y

Answer yes to any questions instead of prompting.

--execute

Execute the notebook prior to export.

--allow-errors

Continue notebook execution even if one of the cells throws an error and include the error message in the cell output (the default behaviour is to abort conversion). This flag is only relevant if '--execute' was specified, too.

--stdin

read a single notebook file from stdin. Write the resulting notebook with default basename 'notebook.\*'

--stdout

Write notebook output to stdout instead of files.

--inplace

Run nbconvert in place, overwriting the existing notebook (only relevant when converting to notebook format)

--clear-output

Clear output of current file and save in place,

overwriting the existing notebook.

`--no-prompt`

Exclude input and output prompts from converted document.

`--no-input`

Exclude input cells and output prompts from converted document.  
This mode is ideal for generating code-free reports.

`--log-level=<Enum> (Application.log_level)`

Default: 30

Choices: (0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR', 'CRITICAL')

Set the log level by value or name.

`--config=<Unicode> (JupyterApp.config_file)`

Default: ''

Full path of a config file.

`--to=<Unicode> (NbConvertApp.export_format)`

Default: 'html'

The export format to be used, either one of the built-in formats

['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'] or a dotted object name that represents the import path for an `Exporter` class

`--template=<Unicode> (TemplateExporter.template_file)`

Default: ''

Name of the template file to use

`--writer=<DottedObjectName> (NbConvertApp.writer_class)`

Default: 'FilesWriter'

Writer class used to write the results of the conversion

`--post=<DottedOrNone> (NbConvertApp.postprocessor_class)`

Default: ''

PostProcessor class used to write the results of the conversion

`--output=<Unicode> (NbConvertApp.output_base)`

Default: ''

overwrite base name use for output files. can only be used when converting one notebook at a time.

`--output-dir=<Unicode> (FilesWriter.build_directory)`

Default: ''

Directory to write output(s) to. Defaults to output to the directory of each notebook. To recover previous default behaviour (outputting to the current working directory) use . as the flag value.

`--reveal-prefix=<Unicode> (SlidesExporter.reveal_url_prefix)`

Default: ''

The URL prefix for reveal.js (version 3.x). This defaults to the reveal CDN, but can be any url pointing to a copy of reveal.js.

For speaker notes to work, this must be a relative path to a local copy of reveal.js: e.g., "reveal.js".

If a relative path is given, it must be a subdirectory of the current directory (from which the server is run).

See the usage documentation

(<https://nbconvert.readthedocs.io/en/latest/usage.html#reveal-js-html-slideshow>) for more details.

`--nbformat=<Enum> (NotebookExporter.nbformat_version)`

Default: 4

Choices: [1, 2, 3, 4]

The nbformat version to write. Use this to downgrade notebooks.

To see all available configurables, use `--help-all`

#### Examples

-----

The simplest way to use nbconvert is

```
> jupyter nbconvert mynotebook.ipynb
```

which will convert mynotebook.ipynb to the default format (probably HTML).

You can specify the export format with `--to`.

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'].

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes 'base', 'article' and 'report'. HTML includes 'basic' and 'full'. You can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template basic mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb
```

```
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

```
[NbConvertApp] WARNING | pattern 'knn_nb.ipynb' matched no files
```

This application is used to convert notebook files (\*.ipynb) to various other formats.

WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.

Options

-----

Arguments that take values are actually convenience aliases to full Configurables, whose aliases are listed on the help line. For more information on full configurables, see '--help-all'.

--debug

set log level to logging.DEBUG (maximize logging output)

--generate-config

generate default config file

-y

Answer yes to any questions instead of prompting.

--execute

Execute the notebook prior to export.

--allow-errors

Continue notebook execution even if one of the cells throws an error and include the error message in the cell output (the default behaviour is to abort conversion). This flag is only relevant if '--execute' was specified, too.

--stdin

read a single notebook file from stdin. Write the resulting notebook with default basename 'notebook.\*'

`--stdout`

Write notebook output to stdout instead of files.

`--inplace`

Run nbconvert in place, overwriting the existing notebook (only relevant when converting to notebook format)

`--clear-output`

Clear output of current file and save in place, overwriting the existing notebook.

`--no-prompt`

Exclude input and output prompts from converted document.

`--no-input`

Exclude input cells and output prompts from converted document. This mode is ideal for generating code-free reports.

`--log-level=<Enum> (Application.log_level)`

Default: 30

Choices: (0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR', 'CRITICAL')

Set the log level by value or name.

`--config=<Unicode> (JupyterApp.config_file)`

Default: ''

Full path of a config file.

`--to=<Unicode> (NbConvertApp.export_format)`

Default: 'html'

The export format to be used, either one of the built-in formats

['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf',

'python', 'rst', 'script', 'slides'] or a dotted object name that represents the import path for an `Exporter` class

`--template=<Unicode> (TemplateExporter.template_file)`

Default: ''

Name of the template file to use

`--writer=<DottedObjectName> (NbConvertApp.writer_class)`

Default: 'FilesWriter'

Writer class used to write the results of the conversion

`--post=<DottedOrNone> (NbConvertApp.postprocessor_class)`

Default: ''

PostProcessor class used to write the results of the conversion

`--output=<Unicode> (NbConvertApp.output_base)`

Default: ''

overwrite base name use for output files. can only be used when converting one notebook at a time.

`--output-dir=<Unicode> (FilesWriter.build_directory)`

Default: ''

Directory to write output(s) to. Defaults to output to the directory of each notebook. To recover previous default behaviour (outputting to the current working directory) use . as the flag value.

`--reveal-prefix=<Unicode> (SlidesExporter.reveal_url_prefix)`

Default: ''

The URL prefix for reveal.js (version 3.x). This defaults to the reveal CDN, but can be any url pointing to a copy of reveal.js.



For speaker notes to work, this must be a relative path to a local copy of reveal.js: e.g., "reveal.js".

If a relative path is given, it must be a subdirectory of the current directory (from which the server is run).

See the usage documentation

(<https://nbconvert.readthedocs.io/en/latest/usage.html#reveal-js-html-slideshow>) for more details.

--nbformat=<Enum> (NotebookExporter.nbformat\_version)

Default: 4

Choices: [1, 2, 3, 4]

The nbformat version to write. Use this to downgrade notebooks.

To see all available configurables, use `--help-all`

Examples

-----

The simplest way to use nbconvert is

```
> jupyter nbconvert mynotebook.ipynb
```

which will convert mynotebook.ipynb to the default format (probably HTML).

You can specify the export format with `--to`.

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'].

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes 'base', 'article' and 'report'. HTML includes 'basic' and 'full'. You can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template basic mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb
```

```
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

```
[NbConvertApp] WARNING | pattern 'knn_nb.ipynb' matched no files
```